Customer Churn Prediction using Machine Learning

Final Internship Project Report

Name: Harsha Teja Gopireddy

Domain: Data Science and Machine Learning

Internship: Upskill Campus

Table of Contents

1. Introduction

2. Objective

3. Dataset Description

4. Exploratory Data Analysis

5. Methodology

6. Models Used

7. Results and Discussion

8. Conclusion

9. Future Scope

10. Tools and Technologies

11. References

1. Introduction

Customer churn prediction is a critical problem in many industries such as telecommunications, banking, e-commerce, and subscription-based services. Customer churn refers to the situation where customers stop using a company's product or service over a certain period of time. Retaining existing customers is often more cost-effective than acquiring new ones, making churn prediction an important area of focus for businesses. By identifying customers who are likely to leave, companies can take proactive measures such as offering personalized discounts, improving customer service, or addressing specific issues to retain them.

With the advancement of data science and machine learning, it has become possible to analyse large volumes of customer data and extract meaningful patterns. Machine learning models can learn from historical data and make predictions about future customer behaviour. In this project, various machine learning algorithms are used to predict customer churn based on features such as tenure, monthly charges, and total charges.

The goal of this project is not only to build predictive models but also to understand the underlying factors that influence customer churn. This helps in making data-driven decisions and improving overall business strategies.

3. Dataset Description

The dataset used in this project contains information about customers and their interaction with a service. It includes several important features that help in understanding customer behavior and predicting whether a customer is likely to churn. The dataset is structured in tabular format, where each row represents an individual customer and each column represents a specific attribute related to that customer.

The key features in the dataset include CustomerID, Gender, Tenure, Monthly Charges, Total Charges, and Churn status. The CustomerID uniquely identifies each

customer. Gender represents whether the customer is male or female. Tenure indicates the number of months the customer has stayed with the company, which is an important factor in analysing loyalty. Monthly Charges refers to the amount the customer pays every month, while Total Charges represents the total amount paid by the customer over time. The target variable, Churn, indicates whether the customer has left the service (Yes) or continues to use it (No).

Before applying machine learning models, the dataset undergoes preprocessing steps such as handling missing values, encoding categorical variables, and formatting data for analysis. These steps ensure that the data is clean, consistent, and suitable for model training. A sample preview of the dataset is shown below to provide a better understanding of its structure.

4. Exploratory Data Analysis

Exploratory Data Analysis (EDA) is a crucial step in any data science project, as it helps in understanding the structure, patterns, and relationships within the dataset before applying machine learning models. In this project, EDA was performed using visualization techniques and statistical summaries to gain insights into customer behaviour and identify factors influencing churn.

Initially, basic statistical analysis was conducted to understand the distribution of each feature, including measures such as mean, median, and standard deviation. This helped in identifying any irregularities or outliers in the data. Histograms were then used to visualize the distribution of numerical features such as tenure, monthly charges, and total charges. These visualizations provided a clear understanding of how the data is spread and whether it follows any specific pattern.

A correlation heatmap was also generated to examine the relationships between different variables in the dataset. This helped in identifying which features are strongly or weakly related to each other and to the target variable (churn). For example, it was observed that customers with higher monthly charges tend to have

a higher likelihood of churn, while customers with longer tenure are more likely to stay with the service.

EDA plays a vital role in selecting relevant features and improving model performance. The insights gained from this analysis guided the data preprocessing and model-building steps in this project.

5. Methodology

The methodology followed in this project involves a systematic approach to building a machine learning model for predicting customer churn. It includes several key steps such as data preprocessing, feature engineering, model training, and evaluation. Each step plays an important role in ensuring the accuracy and reliability of the final model.

The first step is data preprocessing, where the dataset is cleaned and prepared for analysis. This includes handling missing values, removing irrelevant columns such as Customer ID, and ensuring that all data is in the correct format. Categorical variables, such as gender and churn status, are converted into numerical values using label encoding so that they can be used in machine learning algorithms.

Next, the dataset is divided into features (independent variables) and the target variable (dependent variable). The features include attributes like tenure, monthly charges, and total charges, while the target variable is the churn status. The dataset is then split into training and testing sets, typically using an 80:20 ratio, to evaluate model performance on unseen data.

After preprocessing, multiple machine learning models are implemented, including Logistic Regression, Decision Tree, and Random Forest. These models are trained using the training dataset and then used to make predictions on the test dataset.

Finally, the performance of each model is evaluated using metrics such as accuracy, confusion matrix, and classification report. This helps in comparing different models and selecting the most effective one for predicting customer churn.

6. Models Used

In this project, three different machine learning algorithms were implemented to predict customer churn: Logistic Regression, Decision Tree, and Random Forest. These models were selected to compare different approaches to classification and evaluate their performance on the dataset.

Logistic Regression is a simple and widely used classification algorithm that predicts the probability of a binary outcome. In this case, it is used to classify whether a customer will churn or not. It works well with smaller datasets and provides a good baseline for comparison. Logistic Regression is easy to interpret and helps in understanding the relationship between features and the target variable.

Decision Tree is a non-linear model that makes decisions based on a tree-like structure. It splits the dataset into smaller subsets based on feature values, creating a series of decision rules. This model is easy to visualize and interpret, making it useful for understanding how different factors influence customer churn. However, it may overfit the data if not properly controlled.

Random Forest is an ensemble learning method that combines multiple decision trees to improve prediction accuracy and reduce overfitting. It builds several trees using different subsets of the data and averages their predictions. This makes Random Forest more robust and reliable compared to a single decision tree.

By using these three models, the project provides a comprehensive comparison of different machine learning techniques for churn prediction.

7. Results and Discussion

The performance of the machine learning models implemented in this project was evaluated using various metrics such as accuracy, confusion matrix, and

classification report. These metrics provide a comprehensive understanding of how well each model predicts customer churn.

The results show that all three models—Logistic Regression, Decision Tree, and Random Forest—were able to make predictions on the dataset, but their performance varied slightly. Logistic Regression performed well as a baseline model, providing stable and interpretable results. Decision Tree was able to capture non-linear relationships in the data but showed moderate performance due to the small size of the dataset. Random Forest, being an ensemble method, provided more reliable predictions by combining multiple decision trees, although its performance was also limited by the dataset size.

The confusion matrix helped in understanding how many predictions were correct and where the model made errors. It showed the number of true positives, true negatives, false positives, and false negatives, giving deeper insight into model behaviour. The classification report further detailed precision, recall, and F1-score, which are important for evaluating classification performance beyond accuracy.

It is important to note that the dataset used in this project is relatively small, which can affect the stability and generalization of the models. This is also the reason for some warnings observed during evaluation. Despite this limitation, the models demonstrated the ability to learn patterns from the data and make reasonable predictions.

Overall, the results highlight the effectiveness of machine learning techniques in solving customer churn prediction problems.

8. Conclusion

In this project, a machine learning approach was used to predict customer churn based on various customer-related features such as tenure, monthly charges, and total charges. The objective was to analyse customer data, identify patterns, and build predictive models that can help in determining whether a customer is likely to leave a service. Through the implementation of different machine learning

algorithms, the project successfully demonstrated how data-driven techniques can be applied to solve real-world business problems.

Three models—Logistic Regression, Decision Tree, and Random Forest—were developed and evaluated using appropriate performance metrics such as accuracy, confusion matrix, and classification report. Each model provided valuable insights into customer behaviour, with Random Forest showing relatively better robustness compared to the others. The results indicate that machine learning models can effectively identify patterns in customer data and assist organizations in making informed decisions.

Although the dataset used in this project was small, the overall workflow and methodology remain applicable to larger and more complex datasets. The project highlights the importance of data preprocessing, feature selection, and model evaluation in achieving reliable results.

In conclusion, this project reinforces the significance of machine learning in predictive analytics and its potential to help businesses improve customer retention strategies and enhance overall performance.

9. Future Scope

The current project demonstrates the application of basic machine learning techniques for customer churn prediction; however, there is significant scope for further improvement and enhancement. One of the primary areas for future work is the use of a larger and more diverse dataset, which would improve the model's ability to generalize and produce more accurate predictions. A larger dataset would also reduce issues such as class imbalance and improve evaluation metrics.

Another important improvement can be achieved by applying advanced machine learning algorithms such as Gradient Boosting, XG Boost, or Support Vector Machines. These algorithms are more powerful and can capture complex patterns in data more effectively than basic models. Additionally, hyperparameter tuning

techniques such as Grid Search or Random Search can be used to optimize model performance.

Feature engineering is another area with strong potential. By creating new meaningful features or selecting the most relevant ones, the performance of the model can be significantly improved. Incorporating real-time data processing and deploying the model as a web application or API can make the system more practical and useful for real-world business scenarios.

Furthermore, integrating visualization dashboards and automated reporting tools can help organizations better understand customer behaviour and make data-driven decisions. Overall, the project can be extended into a more advanced and scalable solution for customer retention strategies.

10. Tools and Technologies

This project was implemented using a combination of programming languages, libraries, and development tools that are widely used in the field of Data Science and Machine Learning. The primary programming language used is Python, which is known for its simplicity, flexibility, and extensive support for data analysis and machine learning tasks.

Several important Python libraries were utilized throughout the project. Pandas was used for data manipulation and preprocessing, allowing efficient handling of structured datasets. NumPy was used for numerical computations and array operations. For data visualization, Matplotlib and Seaborn were used to create graphs such as histograms and correlation heatmaps, which helped in understanding the data distribution and relationships between variables.

The machine learning models were implemented using the Scikit-learn library, which provides a wide range of algorithms and tools for classification, model training, and evaluation. It was used to build models such as Logistic Regression, Decision Tree, and Random Forest, as well as to calculate evaluation metrics like accuracy and confusion matrix.

The development and execution of the project were carried out using Visual Studio Code (VS Code), which provided an integrated environment for writing, running, and debugging Python code. Overall, these tools and technologies played a crucial role in successfully completing the project and achieving the desired results.

11. References

The successful completion of this project was supported by various learning resources, online documentation, and educational platforms. These resources provided valuable insights into data science concepts, machine learning algorithms, and practical implementation techniques.

The official documentation of Python libraries played a crucial role in understanding their functionalities and usage. The Pandas Documentation was used for data manipulation and preprocessing techniques, while NumPy Documentation helped in performing numerical operations. Visualization techniques were learned and applied using guidance from Matplotlib and Seaborn Documentation, which provided examples for creating graphs and heatmaps.

For implementing machine learning models, the Scikit-learn Documentation was extensively referred to. It provided clear explanations of algorithms such as Logistic Regression, Decision Tree, and Random Forest, along with practical examples of model training and evaluation.

In addition to official documentation, several online learning platforms such as YouTube tutorials, educational blogs, and data science websites were used to understand concepts like exploratory data analysis, feature engineering, and model evaluation. These resources helped in bridging the gap between theoretical knowledge and practical implementation.

Overall, these references were instrumental in guiding the development of this project and enhancing understanding of machine learning techniques.